

Haotian Ma

Madison, WI, USA hma232@wisc.edu <https://atheavenma.github.io/> www.linkedin.com/in/Atheaven-Ma

INTRODUCTION

Computer Science Ph.D. candidate at the University of Wisconsin–Madison, developing methods to understand how learned representations shape model behavior across interventions, combinations, and changing contexts. My work combines theory, method development, controlled interventions, and large-scale evaluation across **mechanistic interpretability**, **LLM reasoning**, **AI safety**, and **AI for scientific discovery**.

EDUCATION

University of Wisconsin–Madison

Ph.D. in Computer Sciences

• Advisor: Prof. Daifeng Wang

Madison, WI

Sep 2021 – Jun 2027

Southern University of Science and Technology

B.S. in Physics

• Advisor: Prof. Hu Xu

Shenzhen, China

Sep 2016 – Jun 2021

RESEARCH

Faithful Representations and Mechanistic Interpretability

Develop methods to determine what neural networks represent and whether those representations support faithful explanations and interventions.

Sparse Autoencoders Do Not Guarantee Local Features

ICLR 2027 Submission

Haotian Ma, Ruqi Yang

- Demonstrated that accurate sparse autoencoder (SAE) reconstruction does not guarantee localized, interpretable features.
- Proved the same SAE can exhibit different input dependencies despite identical representation samples, and validated this across 5 trained SAEs.

Representation Redistribution in Sparse Autoencoders

AAAI 2027 Submission

Ziye Mei, Haotian Ma

- Found that prompt changes preserving the same request can redistribute sparse feature activity, with patterns persisting across unseen prompt variants.
- Validated across 27 SAEs spanning Gemma, Llama, Mistral, GPT, and Pythia, and assessed how they helped predict failures in frozen safety probes.

Quantification and Analysis of Layer-wise and Pixel-wise Information Discarding

ICML 2022 Spotlight

Haotian Ma, Hao Zhang, Fan Zhou, Yingqing Zhang, Quanshi Zhang

- Introduced information entropy measures that compare how neural networks discard input information across layers, models, and individual pixels.
- Derived a theoretical link between selective information discarding and the efficiency of extracting features relevant to the task.

Interpreting CNNs via Decision Trees

CVPR 2019

Quanshi Zhang, Yu Yang, Haotian Ma, Ying Nian Wu

- Converted hidden CNN features into object part concepts and quantified how those parts drive individual predictions.

Interpretable AI for Scientific Discovery

Develop interpretable AI methods that connect model decisions to scientifically meaningful spatial, spectral, and molecular structure.

Spectral-Spatial Interpretation

NeurIPS 2026

Haotian Ma, Ruqi Yang, Philip Townsend

- Pioneered a framework that jointly identifies which spatial locations drive predictions and which spectral channels matter at each location.
- Outperformed Integrated Gradients, DeepSHAP, KernelSHAP, and RISE across 3 datasets in joint faithfulness, spatial localization, and spectral attribution.

SegJointGene: Joint Cell Segmentation and Spatial Gene Prioritization by Information Entropy Guided CNNs

Bioinformatics 2026

Haotian Ma, Daifeng Wang

- Developed an information entropy framework for cell segmentation and gene prioritization, improving molecular signal assignment accuracy by 5–20%.
- Identified neuronal development and synaptic signaling genes in brain tissue and lineage markers in human tonsil.

Spatial Phenotyping Linking Single Cell Genomics to Disease Pathology through Joint Deep Representation Learning

ISMB 2026 Poster

Chenfeng He*, Haotian Ma*, Pubudu Kumarage, Xuerou Li, Kalpana Hanthanan Arachchilage, Shuang Liu, Daifeng Wang

- Learned joint representations of disease phenotype, gene expression, and spatial organization to identify cellular niches associated with disease.
- Integrated single cell RNA sequencing with spatial transcriptomics to transfer phenotype signals across cohorts and map disease pathology.

Compositionality, Intervention, and Systematic Reasoning in LLMs

Study when model components and structural transformations can be combined, reused, or generalized reliably in unseen settings.

Behavioral Compositionality of Sparse Autoencoder Representations

ICLR 2027 Submission

Haotian Ma, Ziye Mei, Frederic Sala

- Characterized downstream interactions among SAE features, revealing combined effects differ from the sum of individual feature effects.
- Pairwise modeling improved combined steering prediction in Llama-3.1 8B by 33%.

Certifying Reuse for LLM Components from Local Interventions

ICLR 2027 Submission

Haotian Ma

- Formulated a framework to certify when locally tested LLM components can be reused in combinations never evaluated directly.
- Proved a few targeted tests can certify component reuse in untested combinations without evaluating every combination.

SymSelect: Learning Which Transformations to Trust for Systematic Reasoning

NAACL 2027 Submission

Haotian Ma, Ruqi Yang

- Established representation similarity as a signal for selecting useful transformations without validity labels.
- Recovered 77 to 92% of supervised alignment gains on MNS and I-RAVEN across ResNet, ViT, CLIP, and Context Choice Transformer.

Reliable LLMs and AI Safety

Study how AI systems remain reliable under uncertain preferences, underspecified user identities, and risks of private information leakage.

The Horizon of Reward Guidance

ICLR 2027 Submission

Haotian Ma

- Defined the support horizon to measure how long existing preference feedback can continue guiding a changing policy.
- Proved this horizon can expire under ordinary policy gradient and showed targeted preference data can extend it.

Gender Attribution by Proxy: Restraint, Fallback, and Rationalization in LLMs

AAAI 2027 Submission

Chenkun Jiang*, Haotian Ma*, Kai Chen, Haiyuan Liu

- Introduced gender attribution by proxy, showing how LLMs turn ordinary user context into unsupported gender assumptions.
- Found that 6 LLMs systematically converted ordinary contextual cues into unsupported gender attributions when user gender was unspecified.

Interpretable Complex-Valued Neural Networks for Privacy Protection

ICLR 2020

Liyao Xiang, Hao Zhang, Haotian Ma, Yifan Zhang, Jie Ren, Quanshi Zhang

- Designed neural networks that encode intermediate features as complex values with randomized phases while preserving accurate predictions.
- Reduced attackers’ ability to reconstruct inputs or infer hidden properties across 8 neural networks.

WORK EXPERIENCE

Waisman Center, University of Wisconsin–Madison

Madison, WI

Graduate Research Assistant. Led methodological development and quantitative analysis for interpretable methods in spatial omics.

Sep 2023 – Present

- Advisor: Prof. Daifeng Wang

Department of Computer Sciences, University of Wisconsin–Madison

Madison, WI

Graduate Research Assistant. Developed symbolic labeling rules and evaluation pipelines for automated weak supervision.

Sep 2021 – Jun 2022

- Advisors: Prof. Aws Albarghouthi and Prof. Frederic Sala

Center for Vision, Cognition, Learning, and Autonomy (VCLA), UCLA

Los Angeles, CA

Research Intern. Explored energy-based modeling and efficient sampling methods for protein discovery.

Jul 2020 – Oct 2020

- Advisor: Prof. Ying Nian Wu

Department of Physics, Southern University of Science and Technology

Shenzhen, China

Research Assistant. Applied generative models and efficient sampling methods to identify stable crystal structures.

Sep 2019 – Jun 2020

- Advisor: Prof. Hu Xu

John Hopcroft Center, Shanghai Jiao Tong University

Shanghai, China

Research Intern. Investigated interpretable methods to explain the prediction logic of convolutional neural networks.

Jul 2018 – Jul 2019

- Advisor: Prof. Quanshi Zhang

SKILLS

- **LLMs & Agent Systems:** LLM Evaluation, LLM Inference, Prompt Engineering, Synthetic Data Generation, Multi-Agent Orchestration, Large-Scale Agent Experimentation, Evaluation Harnesses, vLLM
- **Machine Learning & Interpretability:** Mechanistic Interpretability, Sparse Autoencoders, Representation Analysis, Representation Interventions, LLM Steering, Feature Attribution, Multimodal Learning, Transformers
- **GPU, HPC & Performance Engineering:** CUDA/GPU, GPU Parallelism, Parallel Experiment Execution, Performance Optimization, Slurm/HPC, Linux, Bash, Docker
- **Statistical & Scientific Computing:** Statistical Inference, Hierarchical Bayesian Modeling, Low-Rank Modeling, Image Segmentation, Spatial Omics, Spatial Transcriptomics, Spatial Proteomics
- **Programming & Tools:** Python, C++, R, SQL, PyTorch, TensorFlow, scikit-learn, Git

AWARDS

- **Silver Reviewer**, ICML, 2026
- **Student and Young Professional Award**, ACM BCB, 2025
- **UW–Madison Computer Sciences Scholarship**, 2024
- **UW–Madison Computer Sciences First-Year Scholarship**, 2021
- **Second Prize**, ASC Student Supercomputer Challenge, 2018
- **Silver Medal**, International Genetically Engineered Machine (iGEM) Competition, 2018
- **Outstanding Participant, Ranked 6th of 120**, Nankai University National High School Philosophy Summer Program, 2015
- **National Second Prize**, 6th National Mathematics, Physics and Chemistry Competition for Secondary School Students, 2014

PROFESSIONAL SERVICE

- **Journal Reviewer:** Pattern Recognition (Elsevier) and IEEE Transactions on Neural Networks and Learning Systems (TNNLS)
- **Conference Program Committee:** NeurIPS 2026, ECCV 2026, ICML 2026, CVPR 2026, etc.

TEACHING EXPERIENCE

- **Teaching Assistant**, UW–Madison: Machine Learning (CS 760)
- **Teaching Assistant**, UW–Madison: Introduction to Artificial Intelligence (CS 540)
- **Teaching Assistant**, UW–Madison: Data Science Programming II (CS 320)
- **Education Assistant**, UW–Madison ALP Precollege: Introduction to Neuroscience